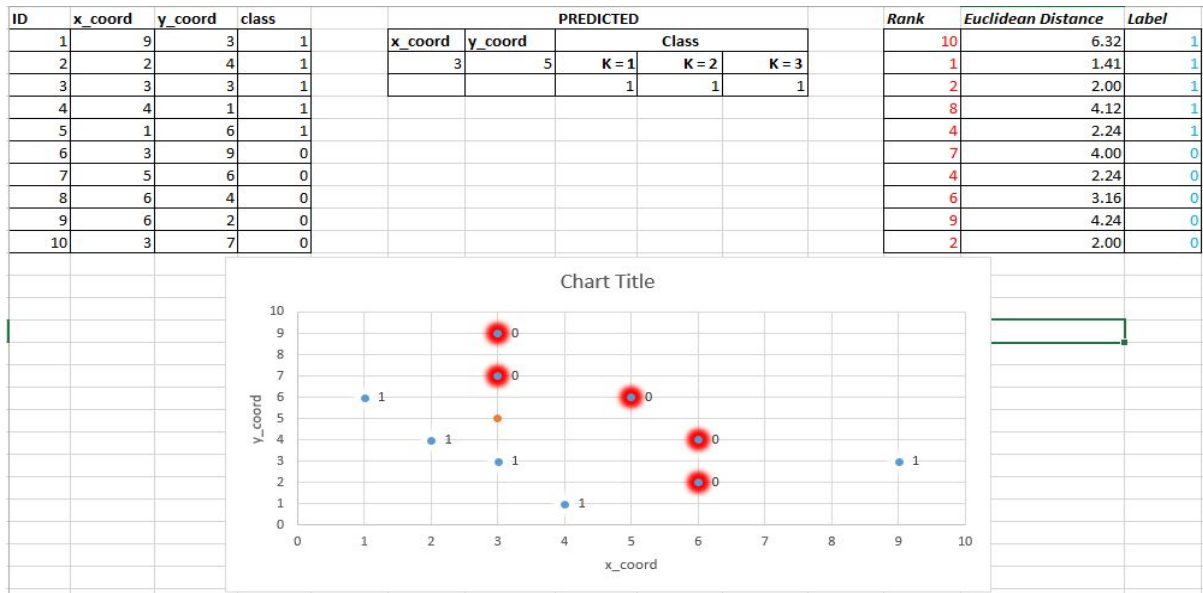


## HOMEWORK 10

### Kenigbolo Meya Stephen

1. Take the following data and simulate K-NN algorithm to predict the class probabilities of points (3, 5) and (4, 6). Report the probabilities with K=1, K=2 and K=3.




As can be seen in the image above, for coordinates (3, 5) the following classes are obtainable:

K = 1    Class 1

K = 2    Class 1

K = 3    Class 1



As can be seen in the image above, for coordinates (4, 6) the following classes are obtainable:

K = 1    Class 0

K = 2    Class 0

K = 3    Class 0

2. In this task we use diabetes dataset to predict diabetes.

- Split the data randomly on 80% of training and 20% for testing.
- Fit the logistic regression on the training set to predict V9.
- Interpret the model. How the Plasma glucose concentration impacts the odds ratio of having diabetes. What about diabetes pedigree function? Which features do not affect (significantly) the risk of having diabetes?
- Now compute Accuracy, Precision, Recall and F1 score on the test set.

For this task I made use of Weka.

- Split the data randomly on 80% of training and 20% for testing.

To split the data into the test set and training set I simple did the following:

- training set:
  - Loaded the full diabetes dataset
  - selected the "RemovePercentage" filter in the preprocess panel
  - set the percentage for the split to 80%
  - applied the filter
  - saved the generated data as a new file
- test set:
  - Reverted the changes to the loaded dataset
  - set the "invertSelection" property to true in the "RemovePercentage" filter
  - applied the filter
  - saved the generated data as new file
- Fit the logistic regression on the training set to predict V9.

I began by uploading the training set and selected "use training set" from the test options in the classify tab. I then selected Logistic regression in the classifier section (After changing V9 from Numeric to Nominal). The output is as follows

**Test mode:evaluate on training data**  
**=== Classifier model (full training set) ===**  
**Logistic Regression with ridge parameter of 1.0E-8**  
**Coefficients...**

	Class
Variable	0
=====	
V1	-0.1217
V2	-0.0329
V3	0.0106
V4	0.0026
V5	0.0009
V6	-0.1026
V7	-1.0648
V8	-0.0096
Intercept	8.5879

**Odds Ratios...**

```

      Class
Variable    0
=====
V1      0.8854
V2      0.9676
V3      1.0106
V4      1.0026
V5      1.0009
V6      0.9024
V7      0.3448
V8      0.9905
Time taken to build model: 0.02 seconds
=== Evaluation on training set ===
=== Summary ===
Correctly Classified Instances   483       78.6645 %
Incorrectly Classified Instances 131       21.3355 %
Kappa statistic                 0.5046
Mean absolute error             0.3046
Root mean squared error         0.3894
Relative absolute error         67.2014 %
Root relative squared error     81.8171 %
Total Number of Instances      614
=== Detailed Accuracy By Class ===

```

	TP Rate	FP Rate	Precision	Recall	F-Measure	ROC Area	Class
	0.893	0.413	0.803	0.893	0.845	0.837	0
	0.587	0.107	0.744	0.587	0.656	0.837	1
Weighted Avg.	0.787	0.307	0.782	0.787	0.78	0.837	

=== Confusion Matrix ===

a b <-- classified as

358 43 | a = 0

88 125 | b = 1

- Now compute Accuracy, Precision, Recall and F1 score on the test set.

To accomplish this I simply selected the option “supplied test set” in the test set options and added the test data from the first step. The output are as follows

```

Test mode:user supplied test set: size unknown (reading incrementally)
=== Classifier model (full training set) ===
Logistic Regression with ridge parameter of 1.0E-8
Coefficients...
      Class
Variable    0
=====
V1      -0.1217
V2      -0.0329
V3       0.0106

```

V4 0.0026  
V5 0.0009  
V6 -0.1026  
V7 -1.0648  
V8 -0.0096  
Intercept 8.5879

**Odds Ratios...**

Variable	Class
	0
V1	0.8854
V2	0.9676
V3	1.0106
V4	1.0026
V5	1.0009
V6	0.9024
V7	0.3448
V8	0.9905

Time taken to build model: 0.02 seconds

=== Evaluation on test set ===

=== Summary ===

Correctly Classified Instances	119	77.2727 %
Incorrectly Classified Instances	35	22.7273 %
Kappa statistic	0.4731	
Mean absolute error	0.3176	
Root mean squared error	0.3973	
Relative absolute error	69.5879 %	
Root relative squared error	82.8984 %	
Total Number of Instances	154	

=== Detailed Accuracy By Class ===

	TP Rate	FP Rate	Precision	Recall	F-Measure	ROC Area	Class
	0.899	0.455	0.781	0.899	0.836	0.848	0
	0.545	0.101	0.75	0.545	0.632	0.848	1
Weighted Avg.	0.773	0.328	0.77	0.773	0.763	0.848	

=== Confusion Matrix ===

a b <-- classified as

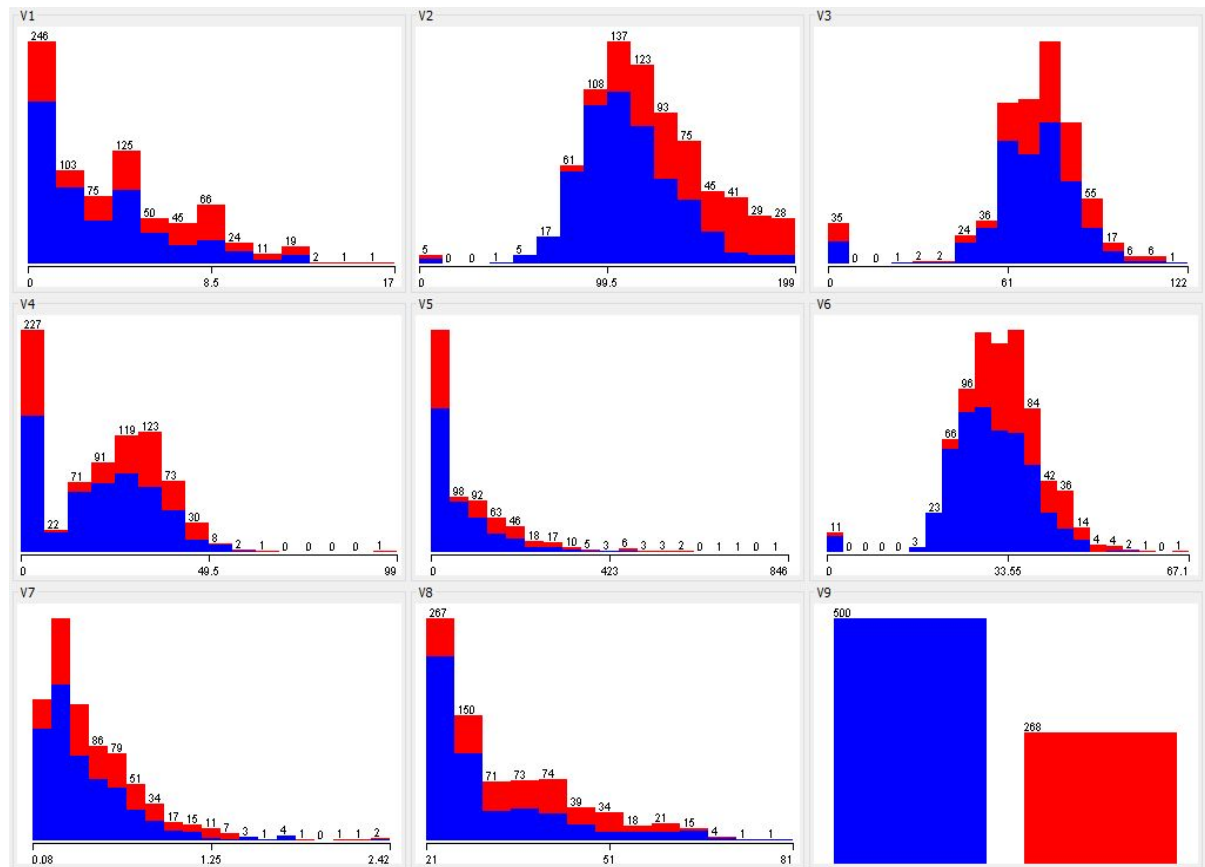
89 10 | a = 0

25 30 | b = 1

- Interpret the model. How the Plasma glucose concentration impacts the odds ratio of having diabetes. What about diabetes pedigree function? Which features do not affect (significantly) the risk of having diabetes?

From the results I can conclude that a high plasma glucose concentration is a strong predictor for future diabetes. Also with a bit of manipulation of removing and adding various attributes I can conclude that the risk of diabetes increases with obesity and older age. The pedigree function does in fact, accurately help estimate the test results for diabetes.

Upon examining the distribution of class values, I noticed that there are 500 negative instances and 268 positive instances. The population is generally young (less than 50 years old) and I believe that some attributes wherein a zero value exists might be due to errors in the data, most especially in mass. Also Plasma glucose concentration, Blood pressure, Skinfold thickness and body mass are normally distributed in the dataset.



### summary of the data in a histogram format

3. Run K-NN on the same data (also using the same setup) to predict diabetes.

- Try different K's (K=1, K=3).
- Report the same scores as before (for each K value).
- Compare the models with F score. Which model has better Accuracy and F score? (logistic or KNN K1 or KNN K3).
- *Optional: plot also roc curves to compare.*

For this task I maintained the train and test data from Task number 2. For obtaining my results I used Weka and implemented the K-NN algorithm in Weka called Ibk.

For K = 1 for the train data the following scores are obtainable

**Test mode:evaluate on training data**

**=== Classifier model (full training set) ===**

**IB1 instance-based classifier**

**using 1 nearest neighbour(s) for classification**

**Time taken to build model: 0 seconds**

**=== Evaluation on training set ===**

**=== Summary ===**

**Correctly Classified Instances      614              100 %**

**Incorrectly Classified Instances      0              0 %**

**Kappa statistic                      1**

**Mean absolute error                  0.0016**

**Root mean squared error              0.0016**

**Relative absolute error               0.3581 %**

**Root relative squared error          0.3411 %**

**Total Number of Instances          614**

**=== Detailed Accuracy By Class ===**

	<b>TP Rate</b>	<b>FP Rate</b>	<b>Precision</b>	<b>Recall</b>	<b>F-Measure</b>	<b>ROC Area</b>	<b>Class</b>
	<b>1</b>	<b>0</b>	<b>1</b>	<b>1</b>	<b>1</b>	<b>1</b>	<b>0</b>
	<b>1</b>	<b>0</b>	<b>1</b>	<b>1</b>	<b>1</b>	<b>1</b>	<b>1</b>
<b>Weighted Avg.</b>	<b>1</b>	<b>0</b>	<b>1</b>	<b>1</b>	<b>1</b>	<b>1</b>	

**=== Confusion Matrix ===**

**a   b   <-- classified as**

**401   0 |   a = 0**

**0   213 |   b = 1**

For K = 1 using test data on model

**Test mode:user supplied test set: size unknown (reading incrementally)**

**=== Classifier model (full training set) ===**

**IB1 instance-based classifier**

**using 1 nearest neighbour(s) for classification**

**Time taken to build model: 0 seconds**

**=== Evaluation on test set ===**

**=== Summary ===**

**Correctly Classified Instances      104              67.5325 %**

**Incorrectly Classified Instances      50              32.4675 %**

**Kappa statistic                      0.3042**

**Mean absolute error                  0.3252**

**Root mean squared error              0.5689**

**Relative absolute error               71.2629 %**

**Root relative squared error          118.7004 %**

**Total Number of Instances          154**

=== Detailed Accuracy By Class ===

	<i>TP Rate</i>	<i>FP Rate</i>	<i>Precision</i>	<i>Recall</i>	<i>F-Measure</i>	<i>ROC Area</i>	<i>Class</i>
	0.727	0.418	0.758	0.727	0.742	0.655	0
	0.582	0.273	0.542	0.582	0.561	0.655	1
<i>Weighted Avg.</i>	0.675	0.366	0.681	0.675	0.678	0.655	

=== Confusion Matrix ===

*a b <-- classified as*

72 27 | *a = 0*

23 32 | *b = 1*

For K = 3 for the train data the following scores are obtainable

*Test mode:evaluate on training data*

=== Classifier model (full training set) ===

*IB1 instance-based classifier*

*using 3 nearest neighbour(s) for classification*

*Time taken to build model: 0 seconds*

=== Evaluation on training set ===

=== Summary ===

<i>Correctly Classified Instances</i>	534	86.9707 %
<i>Incorrectly Classified Instances</i>	80	13.0293 %
<i>Kappa statistic</i>	0.7026	
<i>Mean absolute error</i>	0.1898	
<i>Root mean squared error</i>	0.3035	
<i>Relative absolute error</i>	41.874 %	
<i>Root relative squared error</i>	63.7637 %	
<i>Total Number of Instances</i>	614	

=== Detailed Accuracy By Class ===

	<i>TP Rate</i>	<i>FP Rate</i>	<i>Precision</i>	<i>Recall</i>	<i>F-Measure</i>	<i>ROC Area</i>	<i>Class</i>
	0.938	0.258	0.872	0.938	0.904	0.938	0
	0.742	0.062	0.863	0.742	0.798	0.938	1
<i>Weighted Avg.</i>	0.87	0.19	0.869	0.87	0.867	0.938	

=== Confusion Matrix ===

*a b <-- classified as*

376 25 | *a = 0*

55 158 | *b = 1*

For K = 3 for using the test data on the outcome of the train data the following scores are obtainable

**Test mode:user supplied test set: size unknown (reading incrementally)**

**=== Classifier model (full training set) ===**

**IB1 instance-based classifier**

**using 3 nearest neighbour(s) for classification**

**Time taken to build model: 0 seconds**

**=== Evaluation on test set ===**

**=== Summary ===**

**Correctly Classified Instances 111 72.0779 %**

**Incorrectly Classified Instances 43 27.9221 %**

**Kappa statistic 0.3944**

**Mean absolute error 0.327**

**Root mean squared error 0.472**

**Relative absolute error 71.6536 %**

**Root relative squared error 98.4927 %**

**Total Number of Instances 154**

**=== Detailed Accuracy By Class ===**

	TP Rate	FP Rate	Precision	Recall	F-Measure	ROC Area	Class
	0.778	0.382	0.786	0.778	0.782	0.731	0
	0.618	0.222	0.607	0.618	0.613	0.731	1
Weighted Avg.	0.721	0.325	0.722	0.721	0.721	0.731	

**=== Confusion Matrix ===**

**a b <-- classified as**

**77 22 | a = 0**

**21 34 | b = 1**

### **Conclusion**

From the data above one can comfortably conclude that model where K=3 is more accurate (72.0779%) than the model with K=1(65%) when comparing the number of correctly classified instances for both test data. The F-Score for K=3 is at 0.721 for the test data whereas when that of K = 1 is observed the result is 0.675.

4. In this task we are using diamonds data from the package ggplot2 (data(diamonds)). Build regression models predicting price from the rest of the features, where

A) model 1 has all the features

B) model 2 has all the features + 'carat' and 'depth' of degree 2

C) model 3 has all the features + 3rd degree polynomials of 'carat' and 'depth' (i.e. carat<sup>3</sup>, carat<sup>2</sup>, carat, depth<sup>3</sup>,...)

D) model 4 has all the features + 3rd degree polynomials of 'carat' and 'depth' + 'x','y','z' of degree 2

- in R you can use poly(x,d) to evaluate a polynomial of degree d, e.g. lm(price ~ poly(x,3) + ..., data=diamonds)
- Use the regular 80% train / 20% test split.
- Measure the RMSE for all the models on the train and test set and plot a graph, where on x-axis models are sorted according to the complexity of the model and on y-axis RMSE for train and test split. What do you observe? Can we diagnose under- or overfitting problems?



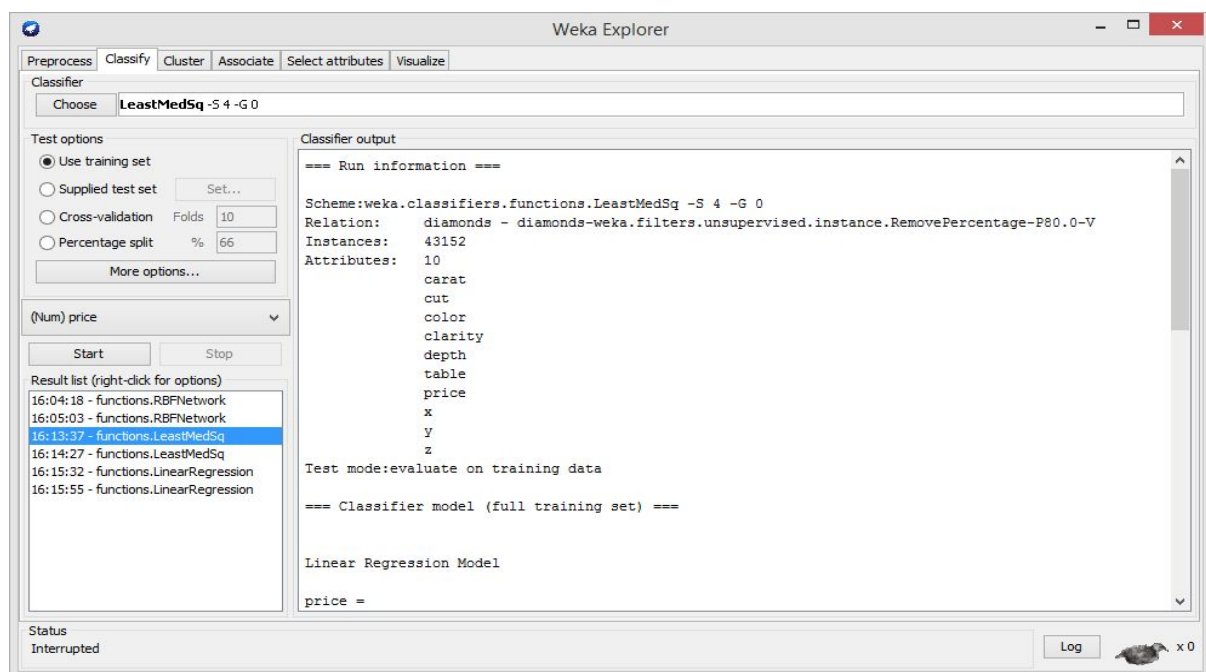
## Solution Task 4

I used both R and Weka for this task. R was used for the pre processing while weka was used for the calculations and graphs. For the first model the data was extracted from R into a csv format where I went ahead to read it in to Weka. In weka I split the data into test and training set and performed the task. The screenshot and output data are as follows

N/B - In order to keep the data meaningful I will only give a summary of the important output data from Weka including the RMSE score. All train/test files attached in their various models in zip file

### Model A (Train)

I used the Least Median Square in Weka to build my model because it implements a better style of linear regression. The summary of the obtainable results for the model are as follows



Test mode:evaluate on training data

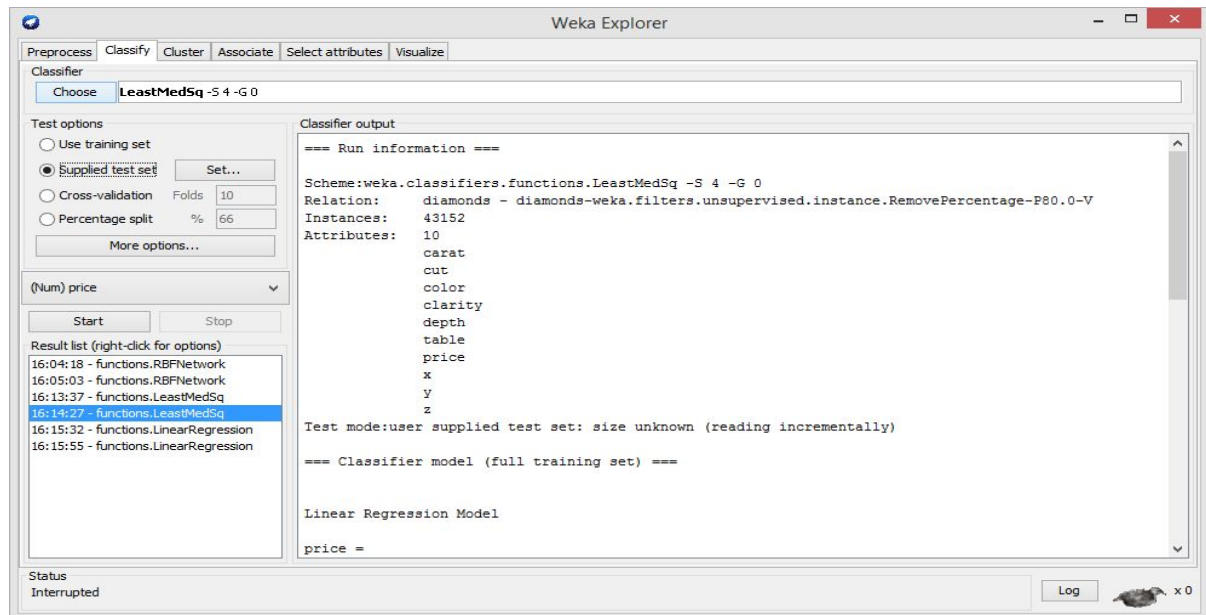
=== Classifier model (full training set) ===

=== Evaluation on training set ===

=== Summary ===

Correlation coefficient	0.946
Mean absolute error	1046.0265
Root mean squared error	2020.2386
Relative absolute error	31.2203 %
Root relative squared error	46.9762 %
Total Number of Instances	43152

### Model A (Test)



Test mode:user supplied test set: size unknown (reading incrementally)

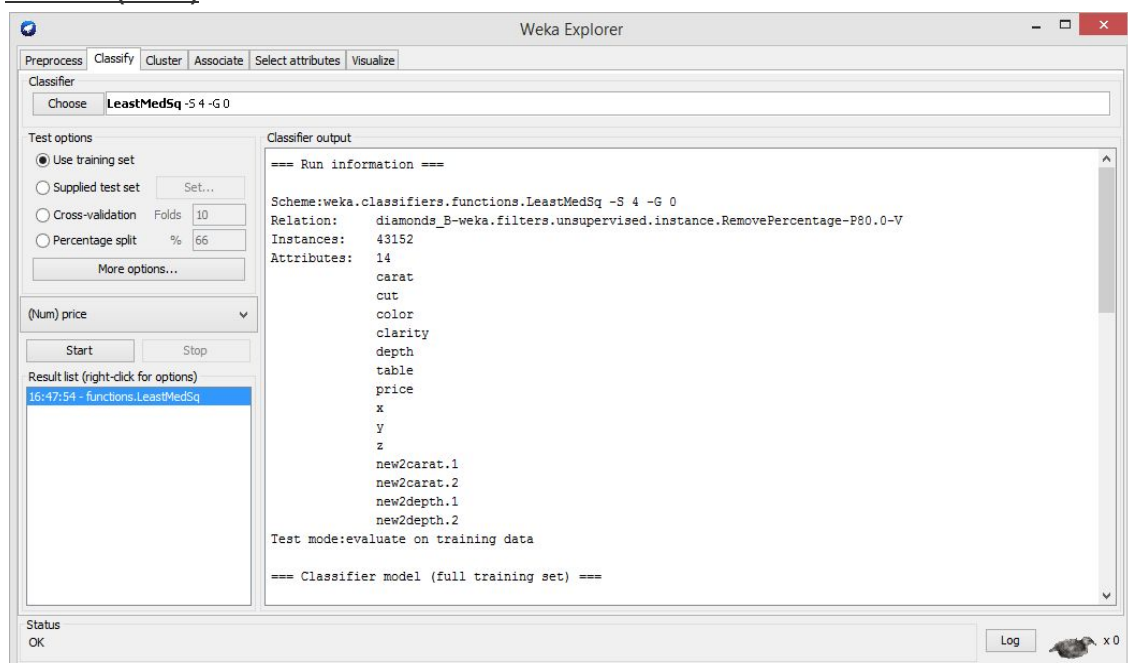
=== Classifier model (full training set) ===

=== Evaluation on test set ===

=== Summary ===

Correlation coefficient	0.9246
Mean absolute error	349.9862
Root mean squared error	424.0939
Relative absolute error	13.6533 %
Root relative squared error	16.1314 %
Total Number of Instances	10788

### Model B (Train)



Test mode:evaluate on training data

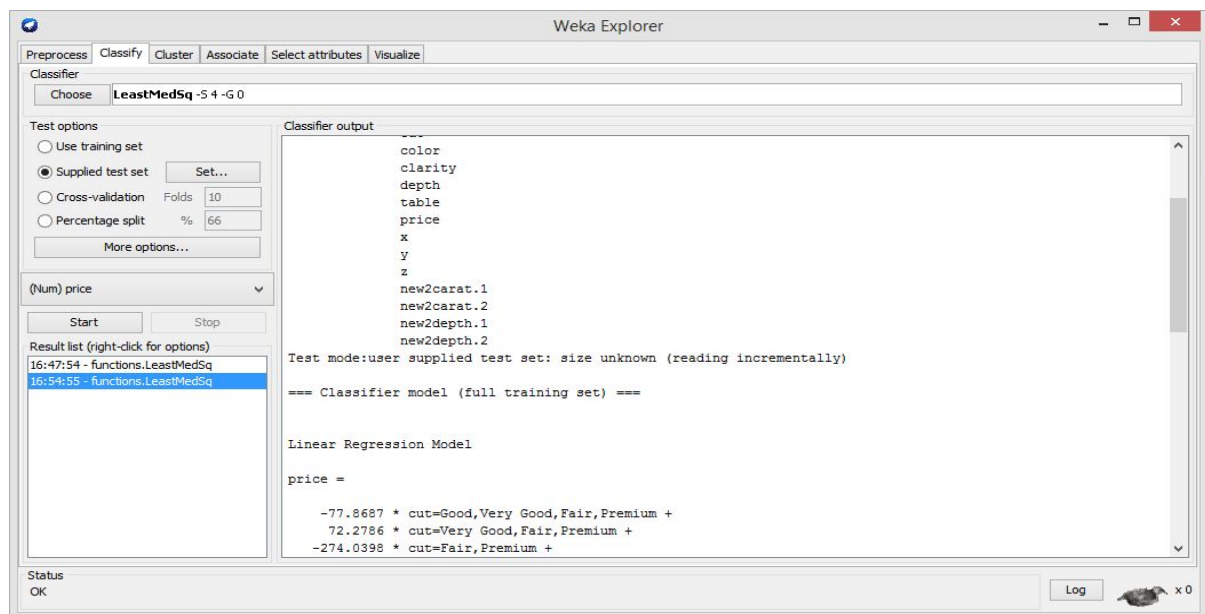
=== Classifier model (full training set) ===

=== Evaluation on training set ===

=== Summary ===

Correlation coefficient	0.9337
Mean absolute error	1141.464
Root mean squared error	2143.0573
Relative absolute error	34.0687 %
Root relative squared error	49.8321 %
Total Number of Instances	43152

### Model B (Test)



Test mode:user supplied test set: size unknown (reading incrementally)

=== Evaluation on test set ===

=== Summary ===

Correlation coefficient	0.8145
Mean absolute error	510.987
Root mean squared error	709.1636
Relative absolute error	19.9341 %
Root relative squared error	26.9746 %
Total Number of Instances	10788

### Summary

For this model I was faced with the task of converting to third degree polynomial which obviously was a challenge to do in Weka so I switched to R, used the formula provided and went ahead to write the data into a csv format and used same on weka for the rest of the Models.

## Model C (Train)

The screenshot shows the Weka Explorer interface with the 'Classify' tab selected. The classifier is set to 'LeastMedSq -S 4 -G 0'. The test options are configured with 'Use training set' selected, 'Cross-validation' set to 10 folds, and 'Percentage split' set to 66%. The result list shows two entries: '17:02:37 - functions.LeastMedSq' and '17:09:12 - functions.LeastMedSq'. The classifier output displays the run information and a list of attributes.

Weka Explorer

Preprocess Classify Cluster Associate Select attributes Visualize

Classifier

Choose LeastMedSq -S 4 -G 0

Test options

☒ Use training set

☐ Supplied test set Set...

☐ Cross-validation Folds 10

☐ Percentage split % 66

More options...

(Num) price

Start Stop

Result list (right-click for options)

17:02:37 - functions.LeastMedSq

17:09:12 - functions.LeastMedSq

Classifier output

=== Run information ===

Scheme:weka.classifiers.functions.LeastMedSq -S 4 -G 0

Relation: diamonds\_C-weka.filters.unsupervised.instance.RemovePercentage-P80.0-V

Instances: 43152

Attributes: 17

carat

cut

color

clarity

depth

table

price

x

y

z

new3carat.1

new3carat.2

new3carat.3

new3depth.1

new3depth.2

new3depth.3

Test mode:evaluate on training data

Status OK Log x0

Test mode:evaluate on training data

=== Evaluation on training set ===

=== Summary ===

Correlation coefficient	0.9096
Mean absolute error	800.9004
Root mean squared error	1810.6231
Relative absolute error	23.9041 %
Root relative squared error	42.1021 %
Total Number of Instances	43152

## Model C (Test)

The screenshot shows the Weka Explorer interface with the 'Classify' tab selected. The classifier is set to 'LeastMedSq -S 4 -G 0'. The test options are configured with 'Supplied test set' selected. The result list shows two entries: '17:02:37 - functions.LeastMedSq' and '17:09:12 - functions.LeastMedSq'. The classifier output displays the run information and a list of attributes.

Weka Explorer

Preprocess Classify Cluster Associate Select attributes Visualize

Classifier

Choose LeastMedSq -S 4 -G 0

Test options

☐ Use training set

☒ Supplied test set Set...

☐ Cross-validation Folds 10

☐ Percentage split % 66

More options...

(Num) price

Start Stop

Result list (right-click for options)

17:02:37 - functions.LeastMedSq

17:09:12 - functions.LeastMedSq

Classifier output

=== Run information ===

Scheme:weka.classifiers.functions.LeastMedSq -S 4 -G 0

Relation: diamonds\_C-weka.filters.unsupervised.instance.RemovePercentage-P80.0-V

Instances: 43152

Attributes: 17

carat

cut

color

clarity

depth

table

price

x

y

z

new3carat.1

new3carat.2

new3carat.3

new3depth.1

new3depth.2

new3depth.3

Test mode:user supplied test set: size unknown (reading incrementally)

Status OK Log x0

Test mode:user supplied test set: size unknown (reading incrementally)

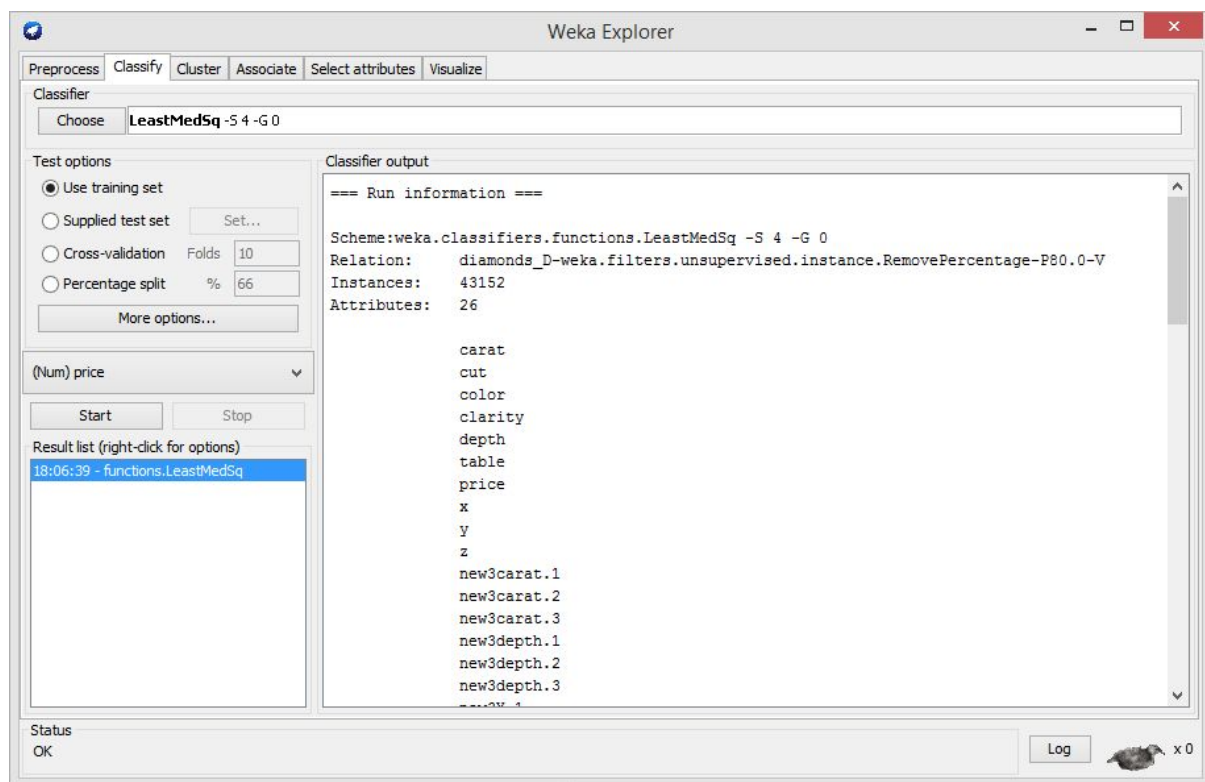
=== Evaluation on test set ===

=== Summary ===

Correlation coefficient	0.7739
Mean absolute error	277.5207
Root mean squared error	535.7505
Relative absolute error	10.8264 %
Root relative squared error	20.3785 %
Total Number of Instances	10788

### Model D (Train)

Like for the previous models I did all the polynomial conversion in R and proceeded to write the data file into csv format which enabled me to preprocess the data with Weka and build a model. The entire flow of calculated attributes can be found in the appendix also.



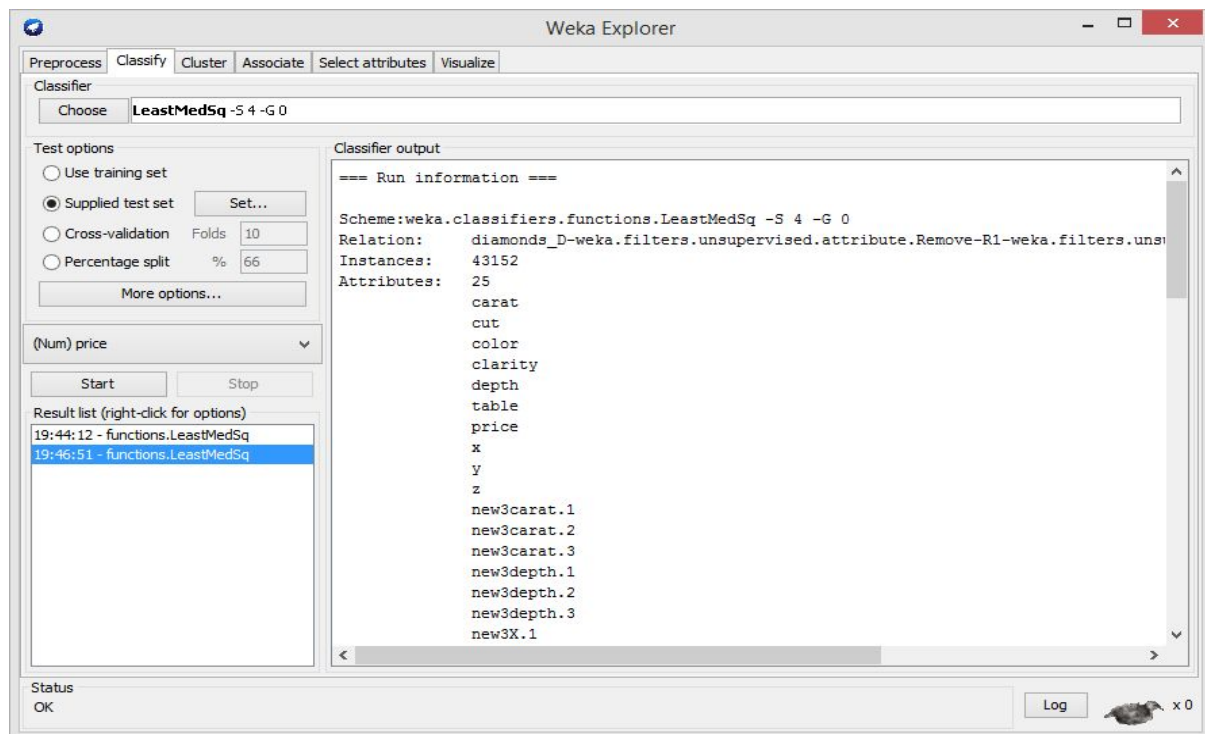
Test mode:evaluate on training data

=== Evaluation on training set ===

=== Summary ===

Correlation coefficient	0.8987
Mean absolute error	1205.845
Root mean squared error	2547.3644
Relative absolute error	35.9903 %
Root relative squared error	59.2333 %
Total Number of Instances	43152

## Model D (Test)




Test mode:user supplied test set: size unknown (reading incrementally)

=== Evaluation on test set ===

=== Summary ===

Correlation coefficient	0.0348
Mean absolute error	491.6804
Root mean squared error	20538.4861
Relative absolute error	19.1809 %
Root relative squared error	781.227 %
Total Number of Instances	10788

## Graph of Models vs RMSE Scores

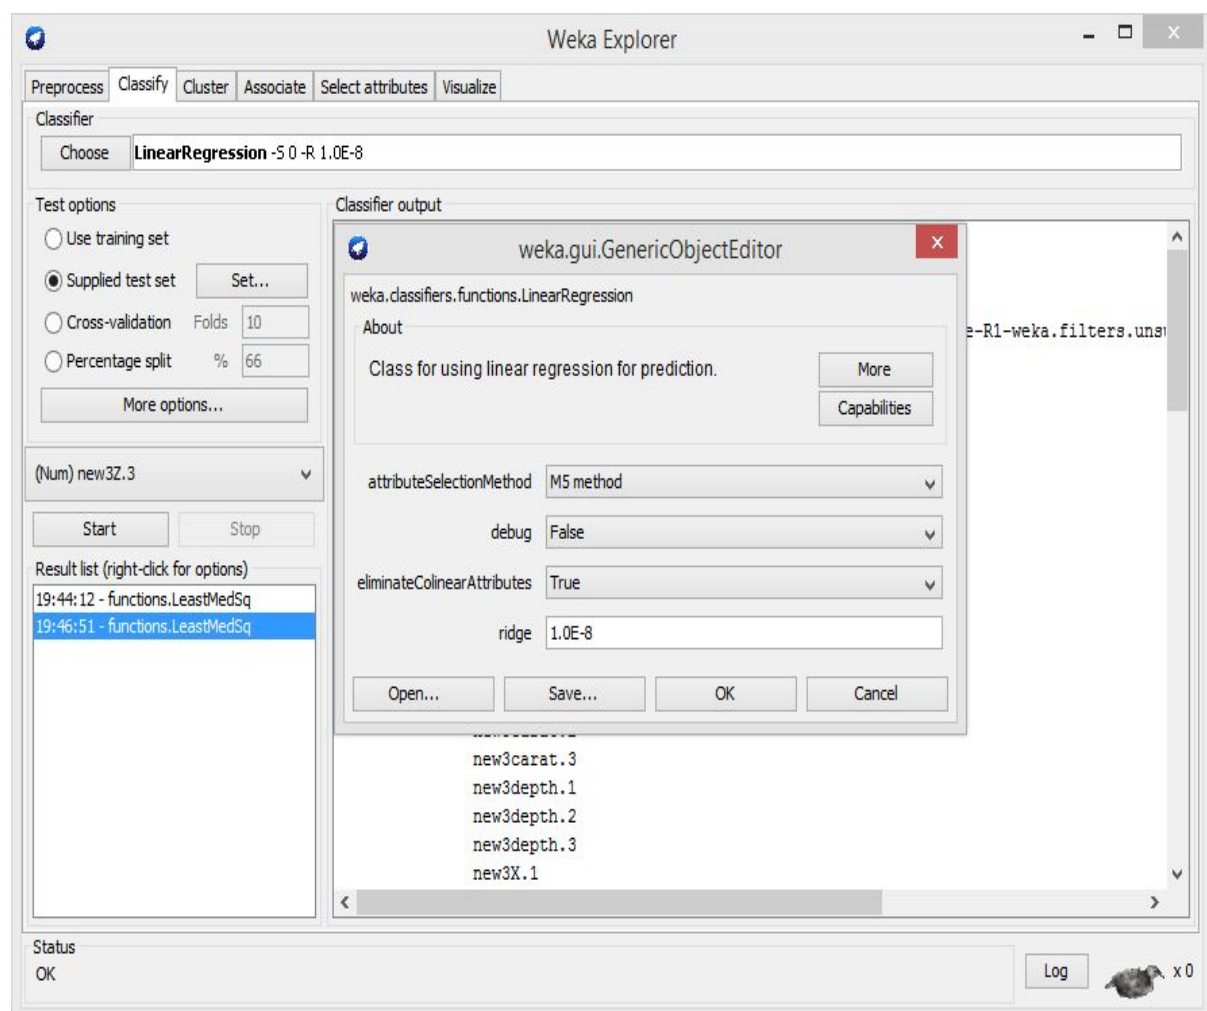


### Observation from Graph

From the Graph above my observation is that an overfitting problem occurs. While both the test and train RMSE for all models between A to C are relatively proportional, there is a disjoint and disproportionate observation between the test and train RMSE score for Model D. This is obviously an overfitting in the model which I believe might be due to some random error or noise instead of the underlying relationship.

7. (optional bonus, 1p). Try ridge and lasso regressions for model 4 from task 4 and add the resulting RMSE of training and test set on the plot generated in task 4. Did it help?

### Model D-Ridge



To perform this Task I made use of the Linear regression function in Weka which implements ridge regression with a default value of 1.0E-8. As can be seen in the screenshot above, this was used to implement ridge regression on Model D. The results for the output of applying Ridge Regression to both Training and Test data on Model D are as follows.



## Model D-Ridge (Train)

The screenshot shows the Weka Explorer interface with the 'Classify' tab selected. The classifier is 'LinearRegression -S 0 -R 1.0E-8'. The 'Test options' section has 'Use training set' selected. The 'Classifier output' pane displays the following data:

```
-135206.8657 * new3X.2 +
25564.4372 * new3X.3 +
743413.3246 * new3Y.1 +
-41064.0521 * new3Y.2 +
-472276.2014 * new3Y.3 +
464343.2192 * new3Z.1 +
3363110.4707 * new3Z.2 +
553207.2084 * new3Z.3 +
6164.0238
```

Time taken to build model: 0.88 seconds

=== Evaluation on training set ===  
=== Summary ===

Correlation coefficient	0.9645
Mean absolute error	775.6269
Root mean squared error	1135.6985
Relative absolute error	23.1498 %
Root relative squared error	26.4082 %
Total Number of Instances	43152

Status: OK

```
=== Evaluation on training set ===
=== Summary ===
Correlation coefficient      0.9645
Mean absolute error        775.6269
Root mean squared error    1135.6985
Relative absolute error     23.1498 %
Root relative squared error 26.4082 %
Total Number of Instances  43152
```

## Model D-Ridge (Test)

The screenshot shows the Weka Explorer interface with the 'Classify' tab selected. The classifier is 'LinearRegression -S 0 -R 1.0E-8'. The 'Test options' section has 'Supplied test set' selected. The 'Classifier output' pane displays the following data:

```
-135206.8657 * new3X.2 +
25564.4372 * new3X.3 +
743413.3246 * new3Y.1 +
-41064.0521 * new3Y.2 +
-472276.2014 * new3Y.3 +
464343.2192 * new3Z.1 +
3363110.4707 * new3Z.2 +
553207.2084 * new3Z.3 +
6164.0238
```

Time taken to build model: 0.92 seconds

=== Evaluation on test set ===  
=== Summary ===

Correlation coefficient	0.0237
Mean absolute error	872.1712
Root mean squared error	33316.1604
Relative absolute error	34.0243 %
Root relative squared error	1267.2542 %
Total Number of Instances	10788

Status: OK



=== Evaluation on test set ===

=== Summary ===

Correlation coefficient	0.0237
Mean absolute error	872.1712
Root mean squared error	33316.1604
Relative absolute error	34.0243 %
Root relative squared error	1267.2542 %
Total Number of Instances	10788

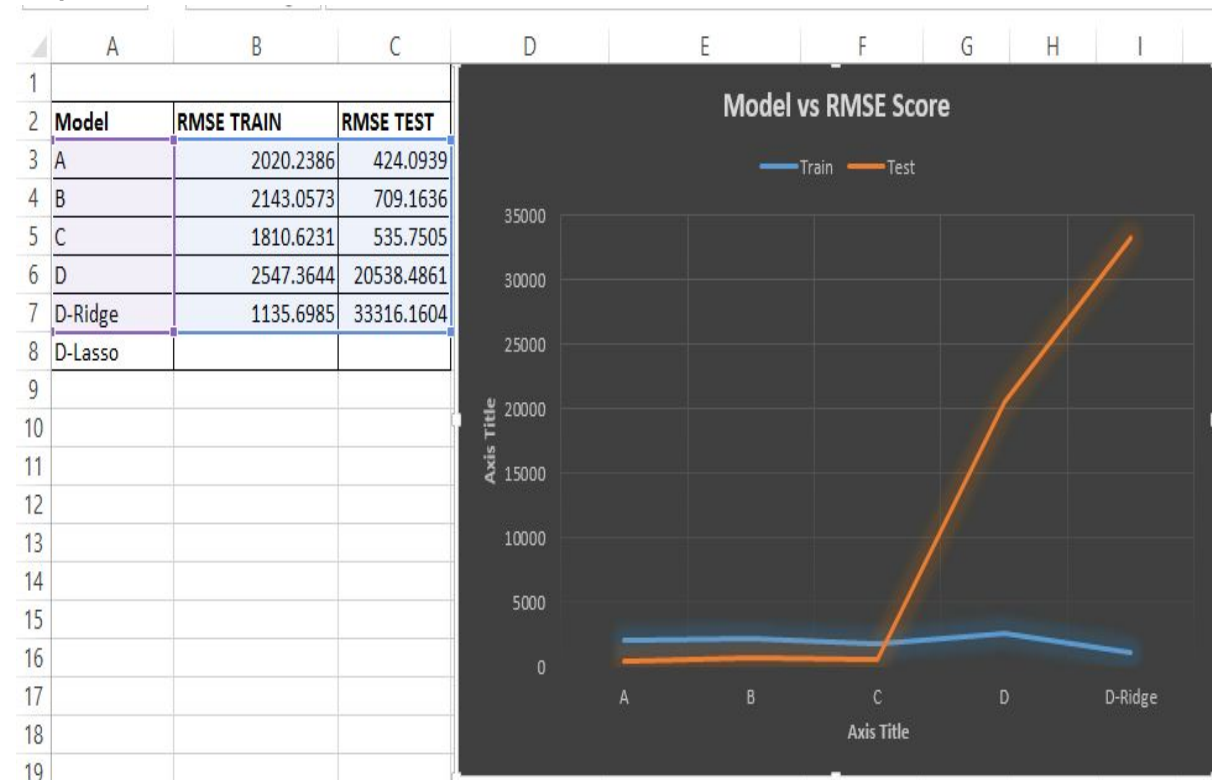
### Model D - Lasso

Using Weka I noticed the absence of any function that implemented Lasso regression hence I did a little research and went through some previous questions asked about lasso regression on Weka. Indeed there is no available in built package for this however there exists an Rplugin that can perform Lasso regression as outlined by the University of Waikato here

<http://weka.8497.n7.nabble.com/Logistic-regression-with-Lasso-regularization-td32836.html>

At the point of this submission I had not successfully implemented the plugin in Weka however I will do that and submit immediately I have been able to implement it.

### Graph of RMSE scores vs Models



### Observation and conclusion

From the screenshot above it is safe to conclude on the basis of applying Ridge regression to Model D that this does not in anyway help us in reducing overfitting to at all as it can be seen that the curve moves upwards as opposed to downwards. There is a slight change in the train set however I personally believe this to be due to the difference in algorithm parameters used in Weka. Unfortunately I cannot compare yet with Lasso regression but I hope to do so soonest

## APPENDIX

### MODEL A TRAIN

Linear Regression Model

price =

5166.6562 \* carat +  
-138.4332 \* cut=Good,Very Good,Fair,Premium +  
89.8772 \* cut=Very Good,Fair,Premium +  
-382.884 \* cut=Fair,Premium +  
396.1977 \* cut=Premium +  
113.5176 \* color=D,F,G,H,I,J +  
-186.174 \* color=F,G,H,I,J +  
-82.5072 \* color=G,H,I,J +  
-66.7112 \* color=H,I,J +  
-262.6146 \* color=I,J +  
-359.4282 \* color=J +  
75.9191 \* clarity=IF,VVS2,VS1,I1,VS2,SI1,SI2 +  
-173.2083 \* clarity=VVS2,VS1,I1,VS2,SI1,SI2 +  
-97.1214 \* clarity=VS1,I1,VS2,SI1,SI2 +  
-2055.9692 \* clarity=I1,VS2,SI1,SI2 +  
1981.4252 \* clarity=VS2,SI1,SI2 +  
-207.7204 \* clarity=SI1,SI2 +  
-449.222 \* clarity=SI2 +  
21.236 \* depth +  
-0.1629 \* table +  
-89.2691 \* x +  
415.9614 \* y +  
15.03 \* z +  
-3416.4567 (Time taken to build model: 32.46 seconds)

### MODEL A TEST

Linear Regression Model

price =

5166.6562 \* carat +  
-138.4332 \* cut=Good,Very Good,Fair,Premium +  
89.8772 \* cut=Very Good,Fair,Premium +  
-382.884 \* cut=Fair,Premium +  
396.1977 \* cut=Premium +  
113.5176 \* color=D,F,G,H,I,J +  
-186.174 \* color=F,G,H,I,J +  
-82.5072 \* color=G,H,I,J +  
-66.7112 \* color=H,I,J +  
-262.6146 \* color=I,J +  
-359.4282 \* color=J +  
75.9191 \* clarity=IF,VVS2,VS1,I1,VS2,SI1,SI2 +  
-173.2083 \* clarity=VVS2,VS1,I1,VS2,SI1,SI2 +  
-97.1214 \* clarity=VS1,I1,VS2,SI1,SI2 +  
-2055.9692 \* clarity=I1,VS2,SI1,SI2 +  
1981.4252 \* clarity=VS2,SI1,SI2 +  
-207.7204 \* clarity=SI1,SI2 +  
-449.222 \* clarity=SI2 +  
21.236 \* depth +  
-0.1629 \* table +  
-89.2691 \* x +  
415.9614 \* y +  
15.03 \* z +

-3416.4567

Time taken to build model: 32.46 seconds

### **MODEL B TRAIN**

Linear Regression Model

price =

-77.8687 \* cut=Good,Very Good,Fair,Premium +  
72.2786 \* cut=Very Good,Fair,Premium +  
-274.0398 \* cut=Fair,Premium +  
252.9161 \* cut=Premium +  
107.5961 \* color=D,F,G,H,I,J +  
-180.8728 \* color=F,G,H,I,J +  
-65.9988 \* color=G,H,I,J +  
-53.3338 \* color=H,I,J +  
-246.579 \* color=I,J +  
-338.7139 \* color=J +  
81.2 \* clarity=IF,VVS2,VS1,I1,VS2,SI1,SI2 +  
-175.5747 \* clarity=VVS2,VS1,I1,VS2,SI1,SI2 +  
-138.0642 \* clarity=VS1,I1,VS2,SI1,SI2 +  
-1999.4805 \* clarity=I1,VS2,SI1,SI2 +  
1954.6267 \* clarity=VS2,SI1,SI2 +  
-236.1905 \* clarity=SI1,SI2 +  
-407.9434 \* clarity=SI2 +  
-203.1043 \* depth +  
19.8103 \* table +  
1109.5151 \* x +  
676.0609 \* y +  
1169.1071 \* z +  
109773.8603 \* new2carat.2 +  
85712.2505 \* new2depth.1 +  
-36251.3557 \* new2depth.2 +  
922.8422

Time taken to build model: 36.31 seconds

### **Model C Train**

Linear Regression Model

price =

-93.0853 \* cut=Good,Very Good,Fair,Premium +  
30.4907 \* cut=Very Good,Fair,Premium +  
-261.7426 \* cut=Fair,Premium +  
276.78 \* cut=Premium +  
122.427 \* color=D,F,G,H,I,J +  
-184.9651 \* color=F,G,H,I,J +  
-36.6564 \* color=G,H,I,J +  
-274.7847 \* color=H,I,J +  
-328.0656 \* color=I,J +  
-176.9376 \* color=J +  
43.3987 \* clarity=IF,VVS2,VS1,I1,VS2,SI1,SI2 +  
-106.1673 \* clarity=VVS2,VS1,I1,VS2,SI1,SI2 +  
4.5306 \* clarity=VS1,I1,VS2,SI1,SI2 +  
-1198.9112 \* clarity=I1,VS2,SI1,SI2 +  
1109.1384 \* clarity=VS2,SI1,SI2 +  
-504.0242 \* clarity=SI1,SI2 +  
-596.9354 \* clarity=SI2 +  
0.4427 \* table +  
637.9665 \* x +

1267.1169 \* y +  
 342.0056 \* z +  
 275136.5102 \* new3carat.1 +  
 -2588.3872 \* new3carat.2 +  
 -295302.4621 \* new3carat.3 +  
 18934.3875 \* new3depth.1 +  
 -21235.1859 \* new3depth.2 +  
 5702.6643 \* new3depth.3 +  
 -7749.3445  
 Time taken to build model: 39.85 seconds

### **Model C Test**

Linear Regression Model

price =  
 -93.0853 \* cut=Good,Very Good,Fair,Premium +  
 30.4907 \* cut=Very Good,Fair,Premium +  
 -261.7426 \* cut=Fair,Premium +  
 276.78 \* cut=Premium +  
 122.427 \* color=D,F,G,H,I,J +  
 -184.9651 \* color=F,G,H,I,J +  
 -36.6564 \* color=G,H,I,J +  
 -274.7847 \* color=H,I,J +  
 -328.0656 \* color=I,J +  
 -176.9376 \* color=J +  
 43.3987 \* clarity=IF,VVS2,VS1,I1,VS2,SI1,SI2 +  
 -106.1673 \* clarity=VVS2,VS1,I1,VS2,SI1,SI2 +  
 4.5306 \* clarity=VS1,I1,VS2,SI1,SI2 +  
 -1198.9112 \* clarity=I1,VS2,SI1,SI2 +  
 1109.1384 \* clarity=VS2,SI1,SI2 +  
 -504.0242 \* clarity=SI1,SI2 +  
 -596.9354 \* clarity=SI2 +  
 0.4427 \* table +  
 637.9665 \* x +  
 1267.1169 \* y +  
 342.0056 \* z +  
 275136.5102 \* new3carat.1 +  
 -2588.3872 \* new3carat.2 +  
 -295302.4621 \* new3carat.3 +  
 18934.3875 \* new3depth.1 +  
 -21235.1859 \* new3depth.2 +  
 5702.6643 \* new3depth.3 +  
 -7749.3445  
 Time taken to build model: 41.55 seconds

### **Model D Train**

Linear Regression Model

price =  
 -0.0001 \* +  
 -99.9338 \* cut=Good,Very Good,Fair,Premium +  
 34.9755 \* cut=Very Good,Fair,Premium +  
 -191.6308 \* cut=Fair,Premium +  
 221.6263 \* cut=Premium +  
 120.4651 \* color=D,F,G,H,I,J +  
 -183.0859 \* color=F,G,H,I,J +  
 -75.8913 \* color=G,H,I,J +  
 -78.2264 \* color=H,I,J +  
 -249.5117 \* color=I,J +  
 -341.3437 \* color=J +

86.7717 \* clarity=IF,VVS2,VS1,I1,VS2,SI1,SI2 +  
 -192.1514 \* clarity=VVS2,VS1,I1,VS2,SI1,SI2 +  
 -59.3148 \* clarity=VS1,I1,VS2,SI1,SI2 +  
 -1625.2994 \* clarity=I1,VS2,SI1,SI2 +  
 1498.2563 \* clarity=VS2,SI1,SI2 +  
 -191.1586 \* clarity=SI1,SI2 +  
 -454.8332 \* clarity=SI2 +  
 711.322 \* depth +  
 -6.2793 \* table +  
 294422.8361 \* new3carat.1 +  
 -159765.937 \* new3carat.2 +  
 -5156.6533 \* new3carat.3 +  
 -243407.892 \* new3depth.1 +  
 -7133.9858 \* new3depth.2 +  
 -2296.6572 \* new3depth.3 +  
 -140196.7508 \* new3X.1 +  
 -270820.3095 \* new3X.2 +  
 -69450.8059 \* new3X.3 +  
 248620.0389 \* new3Y.1 +  
 -470036.477 \* new3Y.3 +  
 374383.4801 \* new3Z.1 +  
 1344389.2257 \* new3Z.2 +  
 179250.2749 \* new3Z.3 +  
 -39817.1002

Time taken to build model: 46.28 seconds

#### Model D Test

Linear Regression Model

price =

-0.0001 \* +  
 -99.9338 \* cut=Good,Very Good,Fair,Premium +  
 34.9755 \* cut=Very Good,Fair,Premium +  
 -191.6308 \* cut=Fair,Premium +  
 221.6263 \* cut=Premium +  
 120.4651 \* color=D,F,G,H,I,J +  
 -183.0859 \* color=F,G,H,I,J +  
 -75.8913 \* color=G,H,I,J +  
 -78.2264 \* color=H,I,J +  
 -249.5117 \* color=I,J +  
 -341.3437 \* color=J +  
 86.7717 \* clarity=IF,VVS2,VS1,I1,VS2,SI1,SI2 +  
 -192.1514 \* clarity=VVS2,VS1,I1,VS2,SI1,SI2 +  
 -59.3148 \* clarity=VS1,I1,VS2,SI1,SI2 +  
 -1625.2994 \* clarity=I1,VS2,SI1,SI2 +  
 1498.2563 \* clarity=VS2,SI1,SI2 +  
 -191.1586 \* clarity=SI1,SI2 +  
 -454.8332 \* clarity=SI2 +  
 711.322 \* depth +  
 -6.2793 \* table +  
 294422.8361 \* new3carat.1 +  
 -159765.937 \* new3carat.2 +  
 -5156.6533 \* new3carat.3 +  
 -243407.892 \* new3depth.1 +  
 -7133.9858 \* new3depth.2 +  
 -2296.6572 \* new3depth.3 +  
 -140196.7508 \* new3X.1 +  
 -270820.3095 \* new3X.2 +  
 -69450.8059 \* new3X.3 +

248620.0389 \* new3Y.1 +  
-470036.477 \* new3Y.3 +  
374383.4801 \* new3Z.1 +  
1344389.2257 \* new3Z.2 +  
179250.2749 \* new3Z.3 +  
-39817.1002  
Time taken to build model: 48.94 seconds